

✓ Roll Number:- 22102B2006

Name:- Khushil Girish Bhimani

Github Link:- <https://github.com/KhushilBhimani2004/Machine-Learning>

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 from sklearn.model_selection import train_test_split
6 from sklearn.linear_model import LinearRegression
7 from sklearn.metrics import mean_squared_error as mse, r2_score
8 from sklearn.metrics import accuracy_score
9 import plotly.express as px
```

```
1 df = pd.read_csv("/content/housing.csv")
```

```
1 df.head(10)
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value
0	-122.23	37.88	41.0	880.0	129.0	322.0	126.0	8.3252	452600.0
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0	1138.0	8.3014	358500.0
2	-122.24	37.85	52.0	1467.0	190.0	496.0	177.0	7.2574	352100.0
3	-122.25	37.85	52.0	1274.0	235.0	558.0	219.0	5.6431	341300.0
4	-122.25	37.85	52.0	1627.0	280.0	565.0	259.0	3.8462	342200.0
5	-122.25	37.85	52.0	919.0	213.0	413.0	193.0	4.0368	269700.0
6	-122.25	37.84	52.0	2535.0	489.0	1094.0	514.0	3.6591	299200.0
7	-122.25	37.84	52.0	3104.0	687.0	1157.0	647.0	3.1200	241400.0
8	-122.26	37.84	42.0	2555.0	665.0	1206.0	595.0	2.0804	226700.0
9	-122.25	37.84	52.0	3549.0	707.0	1551.0	714.0	3.6912	261100.0

Next steps: [Generate code with df](#) [View recommended plots](#)

```
1 # from google.colab import drive
2 # drive.mount('/content/drive')
```

```
1 df.shape
```

```
(20640, 10)
```

```
1 df.columns
```


```
Index(['longitude', 'latitude', 'housing_median_age', 'total_rooms',
      'total_bedrooms', 'population', 'households', 'median_income',
      'median_house_value', 'ocean_proximity'],
      dtype='object')
```

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude             20640 non-null  float64
1   latitude              20640 non-null  float64
2   housing_median_age    20640 non-null  float64
3   total_rooms           20640 non-null  float64
4   total_bedrooms        20433 non-null  float64
5   population            20640 non-null  float64
6   households            20640 non-null  float64
7   median_income         20640 non-null  float64
8   median_house_value    20640 non-null  float64
9   ocean_proximity       20640 non-null  object
dtypes: float64(9), object(1)
```


memory usage: 1.6+ MB

1 df.describe()



	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	populat:
count	20640.000000	20640.000000	20640.000000	20640.000000	20433.000000	20640.000000
mean	-119.569704	35.631861	28.639486	2635.763081	537.870553	1195.697044
std	2.003532	2.135952	12.585558	2181.615252	421.385070	2003.532000
min	-124.350000	32.540000	1.000000	2.000000	1.000000	-124.350000
25%	-121.800000	33.930000	18.000000	1447.750000	296.000000	-121.800000
50%	-118.490000	34.260000	29.000000	2127.000000	435.000000	-118.490000
75%	-118.010000	37.710000	37.000000	3148.000000	647.000000	-118.010000
max	-114.310000	41.950000	52.000000	39320.000000	6445.000000	-114.310000

1 df.isnull().sum()



```

longitude      0
latitude       0
housing_median_age  0
total_rooms    0
total_bedrooms 207
population     0
households     0
median_income  0
median_house_value  0
ocean_proximity  0
dtype: int64

```

1 df.duplicated().sum()

 0


```

1 # med_value = df['total_bedrooms'].median()
2 # med_value

```


1 df =df.dropna(axis=0, how='any')

1 df.sample(10)



	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	populat:
13868	-117.31	34.35	9.0	2404.0	390.0	107
19338	-122.82	38.53	27.0	1823.0	360.0	90
5620	-118.25	33.78	32.0	296.0	139.0	51
4694	-118.37	34.07	52.0	2195.0	435.0	88
2126	-119.71	36.77	11.0	5112.0	1384.0	248
15686	-122.42	37.79	48.0	4506.0	1342.0	198
17892	-121.91	37.36	42.0	3224.0	708.0	194
7441	-118.20	33.94	45.0	1570.0	328.0	132
6629	-118.15	34.16	18.0	1711.0	383.0	147
13279	-117.65	34.10	44.0	1526.0	337.0	83

1 df.isnull().sum()



```

longitude      0
latitude       0
housing_median_age  0
total_rooms    0
total_bedrooms  0
population     0
households     0
median_income  0
median_house_value  0
ocean_proximity  0
dtype: int64

```

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 20433 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude              20433 non-null  float64
1   latitude               20433 non-null  float64
2   housing_median_age     20433 non-null  float64
3   total_rooms            20433 non-null  float64
4   total_bedrooms         20433 non-null  float64
5   population             20433 non-null  float64
6   households             20433 non-null  float64
7   median_income          20433 non-null  float64
8   median_house_value     20433 non-null  float64
9   ocean_proximity        20433 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.7+ MB
```

```
1 unique_val = df['ocean_proximity'].unique()
2 unique_val
```

```
array(['NEAR BAY', '<1H OCEAN', 'INLAND', 'NEAR OCEAN', 'ISLAND'],
      dtype=object)
```

```
1 df['ocean_proximity'] = df['ocean_proximity'].map({'NEAR BAY': 1, '<1H OCEAN': 2, 'INLAND': 3, 'NEAR OCEAN': 4, 'ISLAND': 5})
2 #
3 #
```

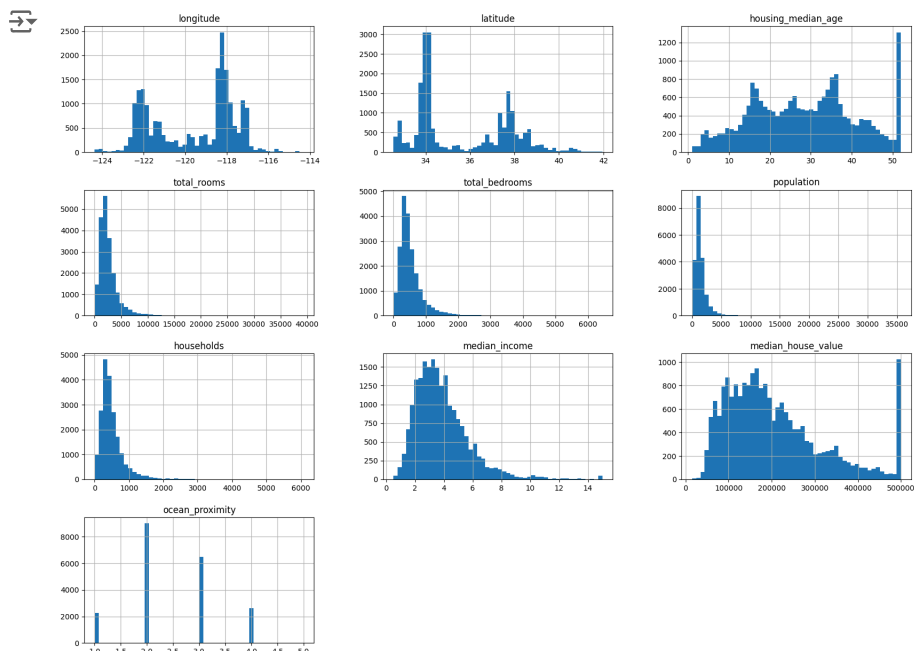
```
1 df.head(10)
```

```
<class 'pandas.core.frame.DataFrame'>
longititude  latitude  housing_median_age  total_rooms  total_bedrooms  population
0   -122.23      37.88              41.0          880.0           129.0          322.0
1   -122.22      37.86              21.0         7099.0           1106.0         2401.0
2   -122.24      37.85              52.0         1467.0           190.0          496.0
3   -122.25      37.85              52.0         1274.0           235.0          558.0
4   -122.25      37.85              52.0         1627.0           280.0          565.0
5   -122.25      37.85              52.0          919.0           213.0          413.0
6   -122.25      37.84              52.0         2535.0           489.0         1094.0
7   -122.25      37.84              52.0         3104.0           687.0         1157.0
8   -122.26      37.84              42.0         2555.0           665.0         1206.0
9   -122.25      37.84              52.0         3549.0           707.0         1551.0
```

Next steps:

[Generate code with df](#)
[View recommended plots](#)

```
1 df.hist(bins=50, figsize=(20,15))
2 plt.show()
```




```
1 df.corr()
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms
longitude	1.000000	-0.924616	-0.109357	0.045480	0.061626
latitude	-0.924616	1.000000	0.011899	-0.036667	-0.045190
housing_median_age	-0.109357	0.011899	1.000000	-0.360628	-0.321772
total_rooms	0.045480	-0.036667	-0.360628	1.000000	0.932845
total_bedrooms	0.061626	-0.045190	-0.321772	0.932845	1.000000
population	0.100270	-0.108997	-0.295787	0.857281	0.871956
households	0.056513	-0.071774	-0.302768	0.918992	0.917697
median_income	-0.015550	-0.079626	-0.118278	0.197882	-0.001913
median_house_value	-0.045398	-0.144638	0.106432	0.133294	0.042907
ocean_proximity	0.181198	-0.067980	-0.206178	0.015917	0.000941

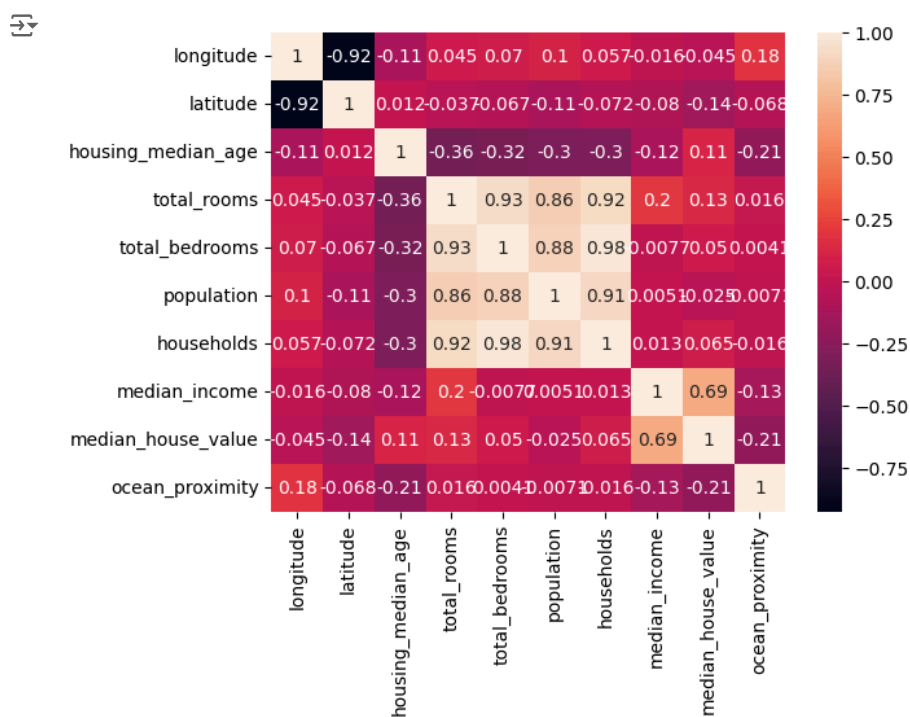
```
1 corrln = df.corr()['median_house_value']
2 round(corrln,2)
```

```
longitude      -0.05
latitude       -0.14
housing_median_age  0.11
total_rooms     0.13
total_bedrooms  0.05
population     -0.03
households      0.06
median_income   0.69
median_house_value 1.00
ocean_proximity -0.21
Name: median_house_value, dtype: float64
```


```
1 df.sample(3000).plot(kind='scatter', x='longitude', y='latitude', s=32, alpha=.8)
2 plt.gca().spines[['top', 'right']].set_visible(False)
```



```
1 ax = sns.heatmap(df.corr(), annot = True)
```



```
1 data = df.head(2000)
2 px.scatter(data, x='median_income', y='median_house_value')
```



```
1 df.columns
2 housing_df = df[['longitude', 'latitude', 'housing_median_age', 'total_rooms',
3                 'total_bedrooms', 'population', 'households', 'median_income',
4                 'ocean_proximity', 'median_house_value']]
```

```
1 housing_df.head()
```



	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population
0	-122.23	37.88	41.0	880.0	129.0	322.0
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0
2	-122.24	37.85	52.0	1467.0	190.0	496.0
3	-122.25	37.85	52.0	1274.0	235.0	558.0
4	-122.25	37.85	52.0	1627.0	280.0	565.0

Next steps:

[Generate code with housing_df](#)
[View recommended plots](#)

```
1 train_pd, test_pd, val_pd = housing_df[:18000], housing_df[18000:19217], housing_df[19215:]
2 len(train_pd), len(test_pd), len(val_pd)
```



```
(18000, 1217, 1218)
```

```
1 X_train, y_train = train_pd.drop('median_house_value', axis=1), train_pd.to_numpy()[:, -1]
2 X_train.head(10)
```



	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population
0	-122.23	37.88	41.0	880.0	129.0	322.0
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0
2	-122.24	37.85	52.0	1467.0	190.0	496.0
3	-122.25	37.85	52.0	1274.0	235.0	558.0
4	-122.25	37.85	52.0	1627.0	280.0	565.0
5	-122.25	37.85	52.0	919.0	213.0	413.0
6	-122.25	37.84	52.0	2535.0	489.0	1094.0
7	-122.25	37.84	52.0	3104.0	687.0	1157.0
8	-122.26	37.84	42.0	2555.0	665.0	1206.0
9	-122.25	37.84	52.0	3549.0	707.0	1551.0

Next steps:

[Generate code with X_train](#)
[View recommended plots](#)

```
1 X_val, y_val = val_pd.to_numpy()[:-1], val_pd.to_numpy()[-1]
2 X_test, y_test = test_pd.to_numpy()[:-1], test_pd.to_numpy()[-1]
```

```
1 X_train.shape, y_train.shape, X_test.shape, y_test.shape, X_val.shape, y_val.shape
```

```
((18000, 9), (18000,), (1217, 9), (1217,), (1218, 9), (1218,))
```

```
1 from sklearn.preprocessing import StandardScaler
2 scaler = StandardScaler()
3 X_train = scaler.fit_transform(X_train)
4 X_test = scaler.transform(X_test)
5 X_val = scaler.transform(X_val)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning:
```

X does not have valid feature names, but StandardScaler was fitted with feature names

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning:
```

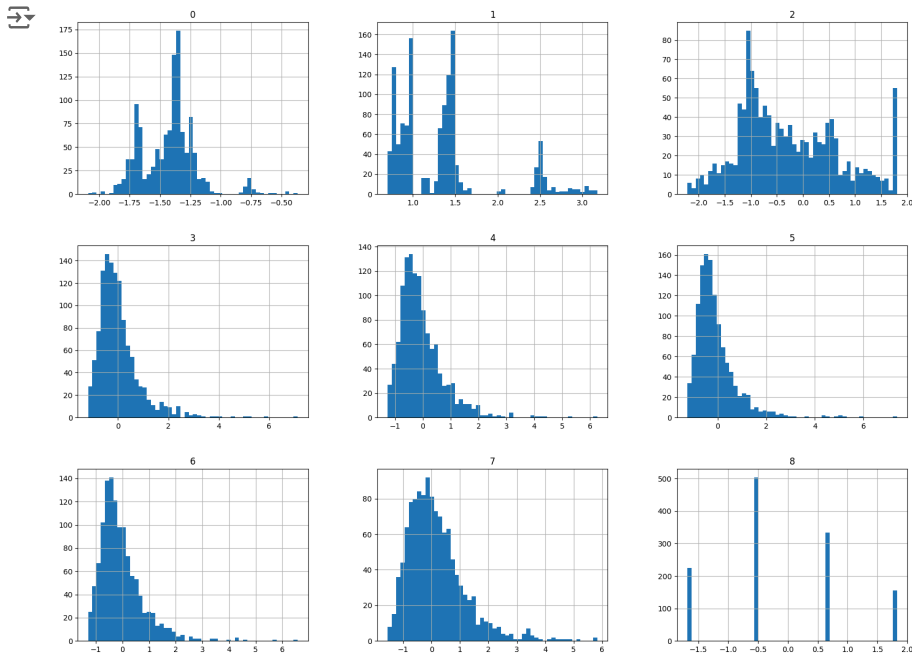
X does not have valid feature names, but StandardScaler was fitted with feature names

```
1 pd.DataFrame(X_train)
```

	0	1	2	3	4	5	6	
0	-1.453822	1.204250	0.935040	-0.795703	-0.964371	-0.970876	-0.972988	2.32936
1	-1.448767	1.194570	-0.642170	2.018961	1.315597	0.843015	1.637958	2.31688
2	-1.458876	1.189729	1.802506	-0.530032	-0.822019	-0.819064	-0.841409	1.76947
3	-1.463931	1.189729	1.802506	-0.617382	-0.717005	-0.764970	-0.733049	0.92303
4	-1.463931	1.189729	1.802506	-0.457617	-0.611991	-0.758863	-0.629850	-0.01915
...
17995	-1.352724	0.947718	-0.799891	0.530841	1.597968	0.432076	1.676658	-0.19260
17996	-1.352724	0.957398	-1.036472	0.345732	1.518624	0.038586	1.392859	-0.15553
17997	-1.347669	0.952558	-0.642170	-0.075631	0.314465	0.044694	0.479544	0.01104
17998	-1.347669	0.952558	-0.326728	-0.255309	-0.362291	-0.141145	-0.309932	0.46045
17999	-1.342614	0.952558	-1.036472	-0.693870	-0.439302	-0.816447	-0.462151	0.16955

18000 rows x 9 columns

```
1 pd.DataFrame(X_test).hist(bins=50, figsize=(20,15))
2 plt.show()
```



```
1 X_train.shape, X_test.shape, X_val.shape,
```

```
((18000, 9), (1217, 9), (1218, 9))
```

```
1 ##Linear Regression Model
```

```
1 # Preprocessing - scaling the data
```

```
2 scaler = StandardScaler()
```

```
3 X_train_scaled = scaler.fit_transform(X_train)
```

```
4 X_val_scaled = scaler.transform(X_val)
```

```
5
```

```
6 # Train the model
```

```
7 lm = LinearRegression().fit(X_train_scaled, y_train)
```

```
1 y_train_pred = lm.predict(X_train_scaled)
```

```
2 y_val_pred = lm.predict(X_val_scaled)
```

```
1 mse_train = mse(y_train, y_train_pred)
```

```
2 rmse_train = mse(y_train, y_train_pred, squared=False)
```

```
3
```

```
4 # Calculate MSE and RMSE for validation set
```

```
5 mse_val = mse(y_val, y_val_pred)
```

```
6 rmse_val = mse(y_val, y_val_pred, squared=False)
```

```
7
```

```
8 # Calculate R2 score for training set
```

```
9 r2_train = r2_score(y_train, y_train_pred)
```

```
10
```

```
11 # Calculate R2 score for validation set
```

```
12 r2_val = r2_score(y_val, y_val_pred)
```



```

1 print(f'Training MSE: {mse_train}')
2 print(f'Training RMSE: {rmse_train}')
3 print(f'Training R²: {r2_train}')

```

Training MSE: 4985623211.241477

```

1 print(f'Validation MSE: {mse_val}')
2 print(f'Validation RMSE: {rmse_val}')
3 print(f'Validation R²: {r2_val}')

```

Validation MSE: 3021634923.4105325
 Validation RMSE: 54969.399882212034
 Validation R²: 0.6626316715336671

Double-click (or enter) to edit


```
1 print('Some predictions on the validation set:', y_val_pred[:5])
```

Some predictions on the validation set: [104479.44966118 192574.74579031 146838.49904684 131771.43881713 109040.18165488]

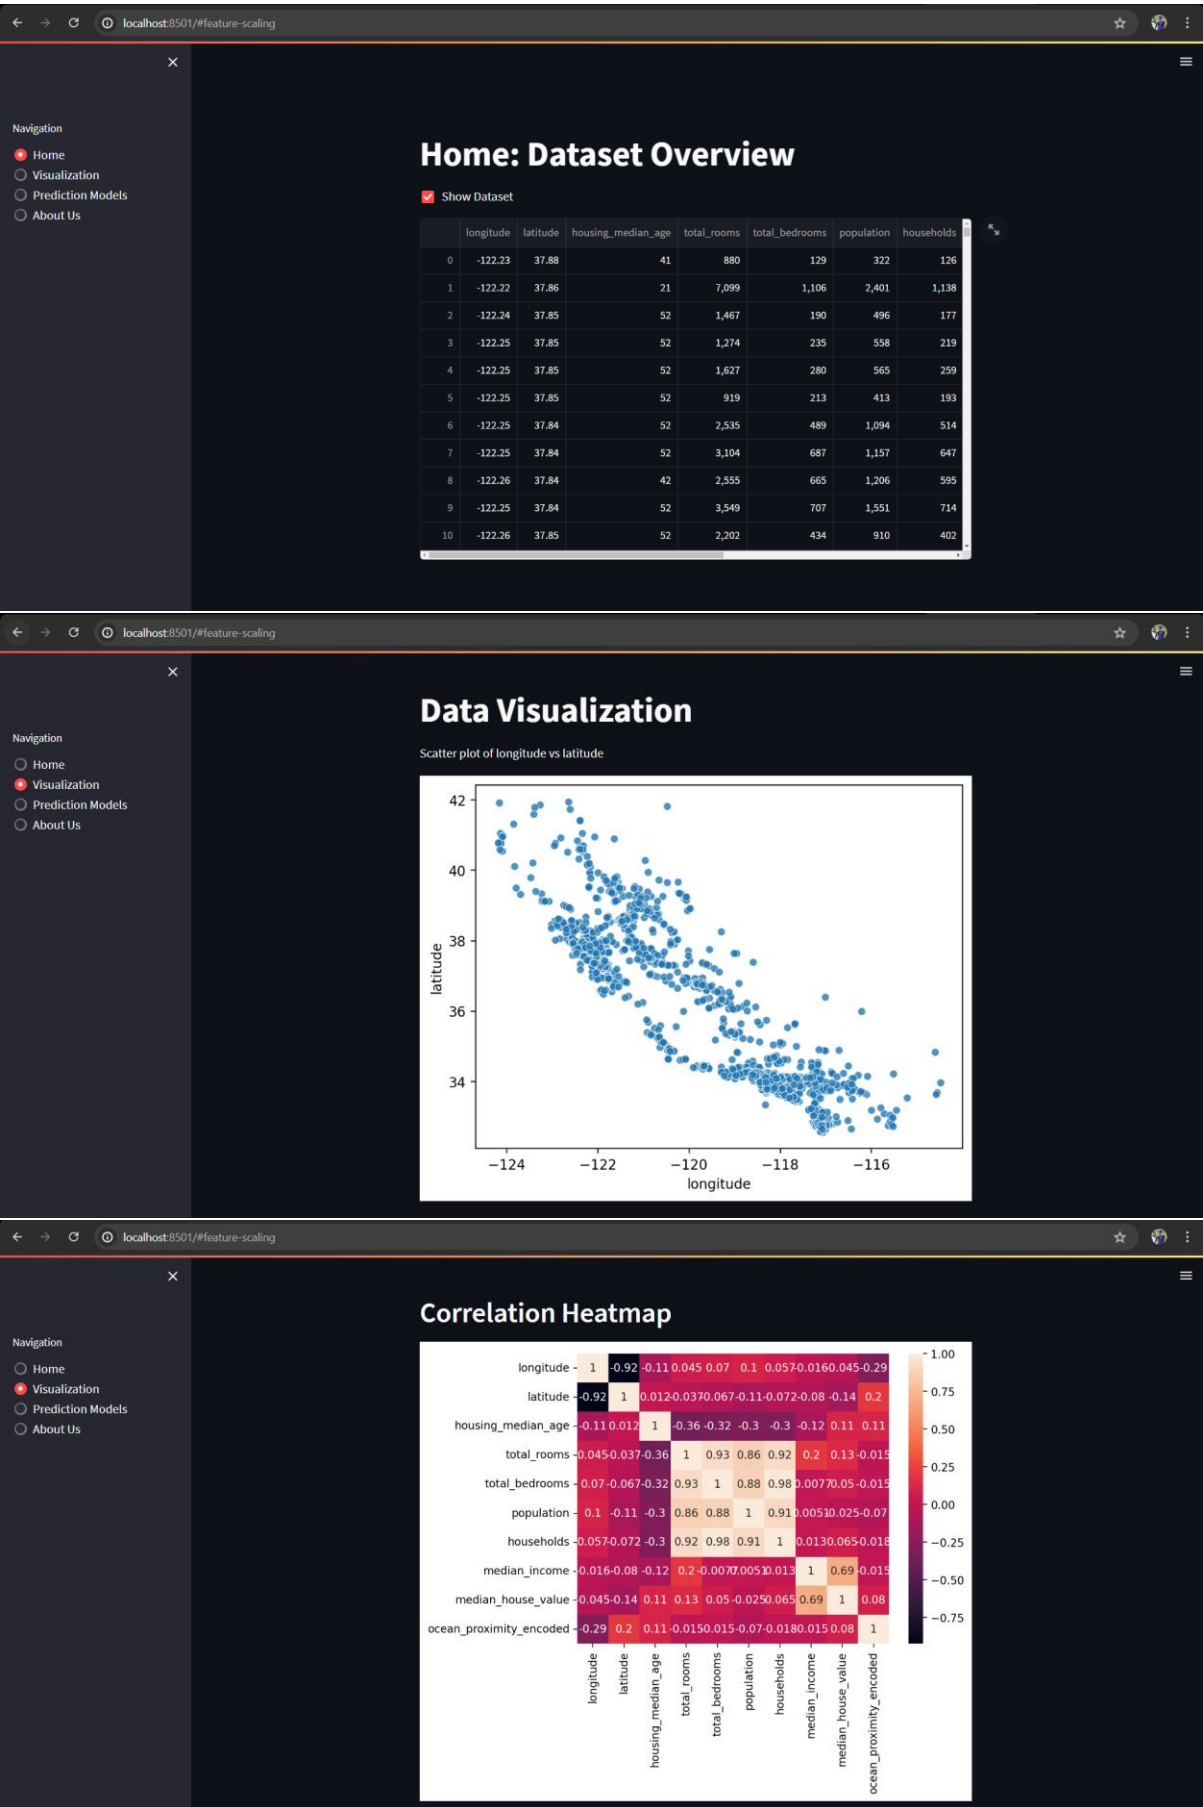
```

1 plt.figure(figsize=(10, 5))
2
3 # Plot actual values
4 plt.scatter(range(len(y_val)), y_val, color='blue', alpha=0.5, label='Actual')
5
6 # Plot predicted values
7 plt.scatter(range(len(y_val)), y_val_pred, color='red', alpha=0.5, label='Predicted')
8
9 plt.xlabel('House Price')
10 plt.ylabel('Value')
11 plt.title('Actual vs Predicted Values')
12 plt.legend()
13 plt.show()
14

```



User Interface -



User Interface -

